

# QUDA-DPC++ migration

Alexei Strelchenko

Scientific Computing Division @ Fermilab

February 26, 2020

# The QUDA library for Lattice Quantum Field Theory

- <http://lattice.github.com/quda> (open source, BSD license)
- GPU backend for BQCD, Chroma, CPS, MILC, TIFR, etc.
- Provides:
  - ▶ Various solvers for all major fermionic discretizations
  - ▶ Additional performance-critical routines needed for gauge field generation
- Maximize performance:
  - ▶ Exploit physical symmetries to minimize memory traffic
  - ▶ Mixed-precision methods
  - ▶ Autotuning for high performance on all CUDA-capable architectures
  - ▶ Domain-decomposed (Schwarz) and multigrid preconditioners for strong scaling
- Objectives:
  - ▶ Prepare QUDA for Intel exascale architectures with "reasonable" efforts

# The QUDA library: autotuning

- QUDA includes an autotuner for ensuring optimal kernel performance
  - ▶ virtual C++ class "Tunable" that is derived for each kernel one wants to autotune
- By default Tunable classes will autotune 1-d CTA size, shared memory size, grid size
  - ▶ Derived specializations do 2-d and 3-d CTA tuning
- Tuned parameters are stored in a `std::map` and dumped to disk for later reuse
- Built in performance metrics and profiling

# QUDA porting efforts

- QUDA major challenges
  - ▶ Low-level optimized compute kernels
  - ▶ (though most of them have x86 counterparts)
  - ▶ CUDA API calls for memory management etc.
- HW/SW specifics:
  - ▶ Some kernels exploit TC (in future releases)
  - ▶ All communication critical kernels will exploit NVSHMEM features

# QUDA porting strategy

- Employ DPC++ CT for source code migration
  - ▶ Option 1: convert CUDA-based versions
  - ▶ Option 2: port x86 versions
- Identify minimal but representative subset of QUDA
- Migration procedure:
  - ▶ Create a compilation Database file
  - ▶ Use Compilation Database File to convert source to DPC++
  - ▶ Verify the Source for Correctness

# Minimal source subset example

```
3 set (QUDA_OBJS
4   # cmake-format: sortable
5   timer.cpp malloc.cpp
6   interface_quda.cpp util_quda.cpp
7   color_spinor_field.cpp color_spinor_util.cu
8   cpu_color_spinor_field.cpp cuda_color_spinor_field.cpp
9   lattice_field.cpp
10  tune.cpp
11  blas_quda.cu reduce_quda.cu
12  random.cu
13  comm_common.cpp ${COMM_OBJS} ${NUMA_AFFINITY_OBJS} ${QIO_UTIL}
14  copy_color_spinor.cu spinor_noise.cu
15  copy_color_spinor_dd.cu copy_color_spinor_ds.cu
16  copy_color_spinor_dh.cu copy_color_spinor_dq.cu
17  copy_color_spinor_ss.cu copy_color_spinor_sd.cu
18  copy_color_spinor_sh.cu copy_color_spinor_sq.cu
19  copy_color_spinor_hd.cu copy_color_spinor_hs.cu
20  copy_color_spinor_hh.cu copy_color_spinor_hq.cu
21  copy_color_spinor_qd.cu copy_color_spinor_qs.cu
22  copy_color_spinor_qh.cu copy_color_spinor_qq.cu
23  quda_cuda_api.cpp
24  version.cpp )
25 # cmake-format: on
26
```

Figure: 1: Source subset

# QUDA porting efforts

- Recent DPCT (beta4)
  - ▶ crashed with the latest version of QUDA
  - ▶ as the first case used with older version

# Conversion script example

```
1 source /opt/oneapi/inteloneapi/dpcpp-ct/latest/env/vars.sh
2
3 QUDA_HOME=/home/astrel/Work/OneAPI/quda-oneapi-v2
4 PROJ_DIR=$QUDA_HOME/lib
5
6 OUTDIR=dpct_output_$(date +%s)
7
8 # copy_color_spinor_d*.cu
9
10 dpct -p -in-root=$PROJ_DIR -out-root=$OUTDIR \
11     --keep-original-code \
12     --extra-arg="-I$QUDA_HOME/include" \
13     --extra-arg="-I./include" \
14     --extra-arg="-std=c++14" \
15     --extra-arg="-I$QUDA_HOME/lib/copy_color_spinor.cuh" \
16     --extra-arg="-I$QUDA_HOME/include/externals" \
17     $QUDA_HOME/lib/copy_color_spinor_dd.cu
18
```

Figure 2: Source subset



- DPCT ignores header files
  - ▶ some \*.h files contain CUDA specific stuff!
- Solution:
  - ▶ Change file extensions, i.e., \*.h  $\rightarrow$  \*.cu
  - ▶ Warning: changing with \*.cuh does not help!
  - ▶ apply dpct to \*.cu files
  - ▶ inspect for correctness and change \*.cu  $\rightarrow$  \*.h

# DPCT issues

- DPCT may not correctly process macros or even fails to generate a dpcpp file

```
115 / DPCT_ORIG      printfQuda("%e Failures: %d / %d = %e\n",  
116 * pow(10.0, -(f+1)/(double)tol), /  
117 /* Complete FAIL! */  
118     printfQuda(  
119         "%e Failures: %d / %d = %e\n",  
120         pow(printfQuda("%e Failures: %d / %d = %e\n",  
121             pow(10.0, -(f + 1) / (double)tol), fail[f],  
122             u.Volume() * N, fail[f] / (double)(u.Volume() * N)),  
123             -(f + 1) / (double)tol),  
124             fail[f], u.Volume() * N, fail[f] / (double)(u.Volume() * N));  
125  
126     }  
127
```

Figure: 2: Macro pre-processing

- DPCT may fail to convert standard CUDA API calls

```
463 / DPCT_ORIG    cudaMemsetAsync(v, 0, bytes, streams[Nstream-1]);  
464             streams[Nstream - 1]->memset(v, 0, bytes);
```

Figure: 2: CUDA API call

- DPCT first experience
  - ▶ Easy to use out of box
  - ▶ Requires an extra inspections of generated codes and still fails on some source files
  - ▶ Requires (manual) preparations for some files
- Deployment:  
<https://github.com/lattice/quda/tree/experimental/oneapi-quda>